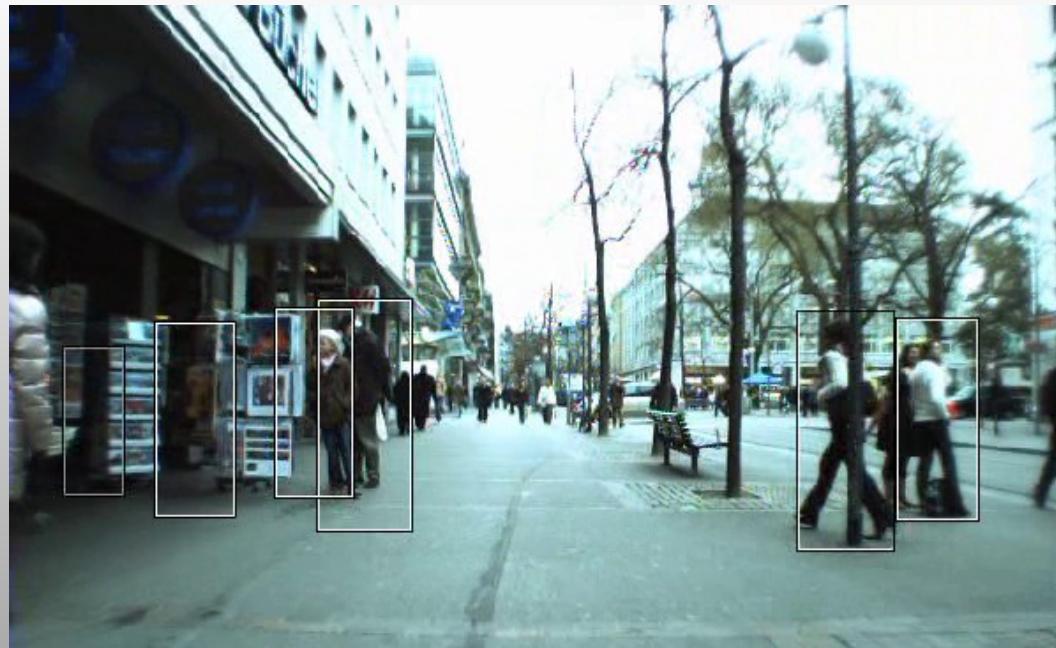


# Fast stixel computation for fast pedestrian detection

R. Benenson, M. Mathias, R. Timofte and L. Van Gool



We want *very* fast  
pedestrian detections



In this paper,  
we reach *high quality*  
pedestrian detection at 165 Hz

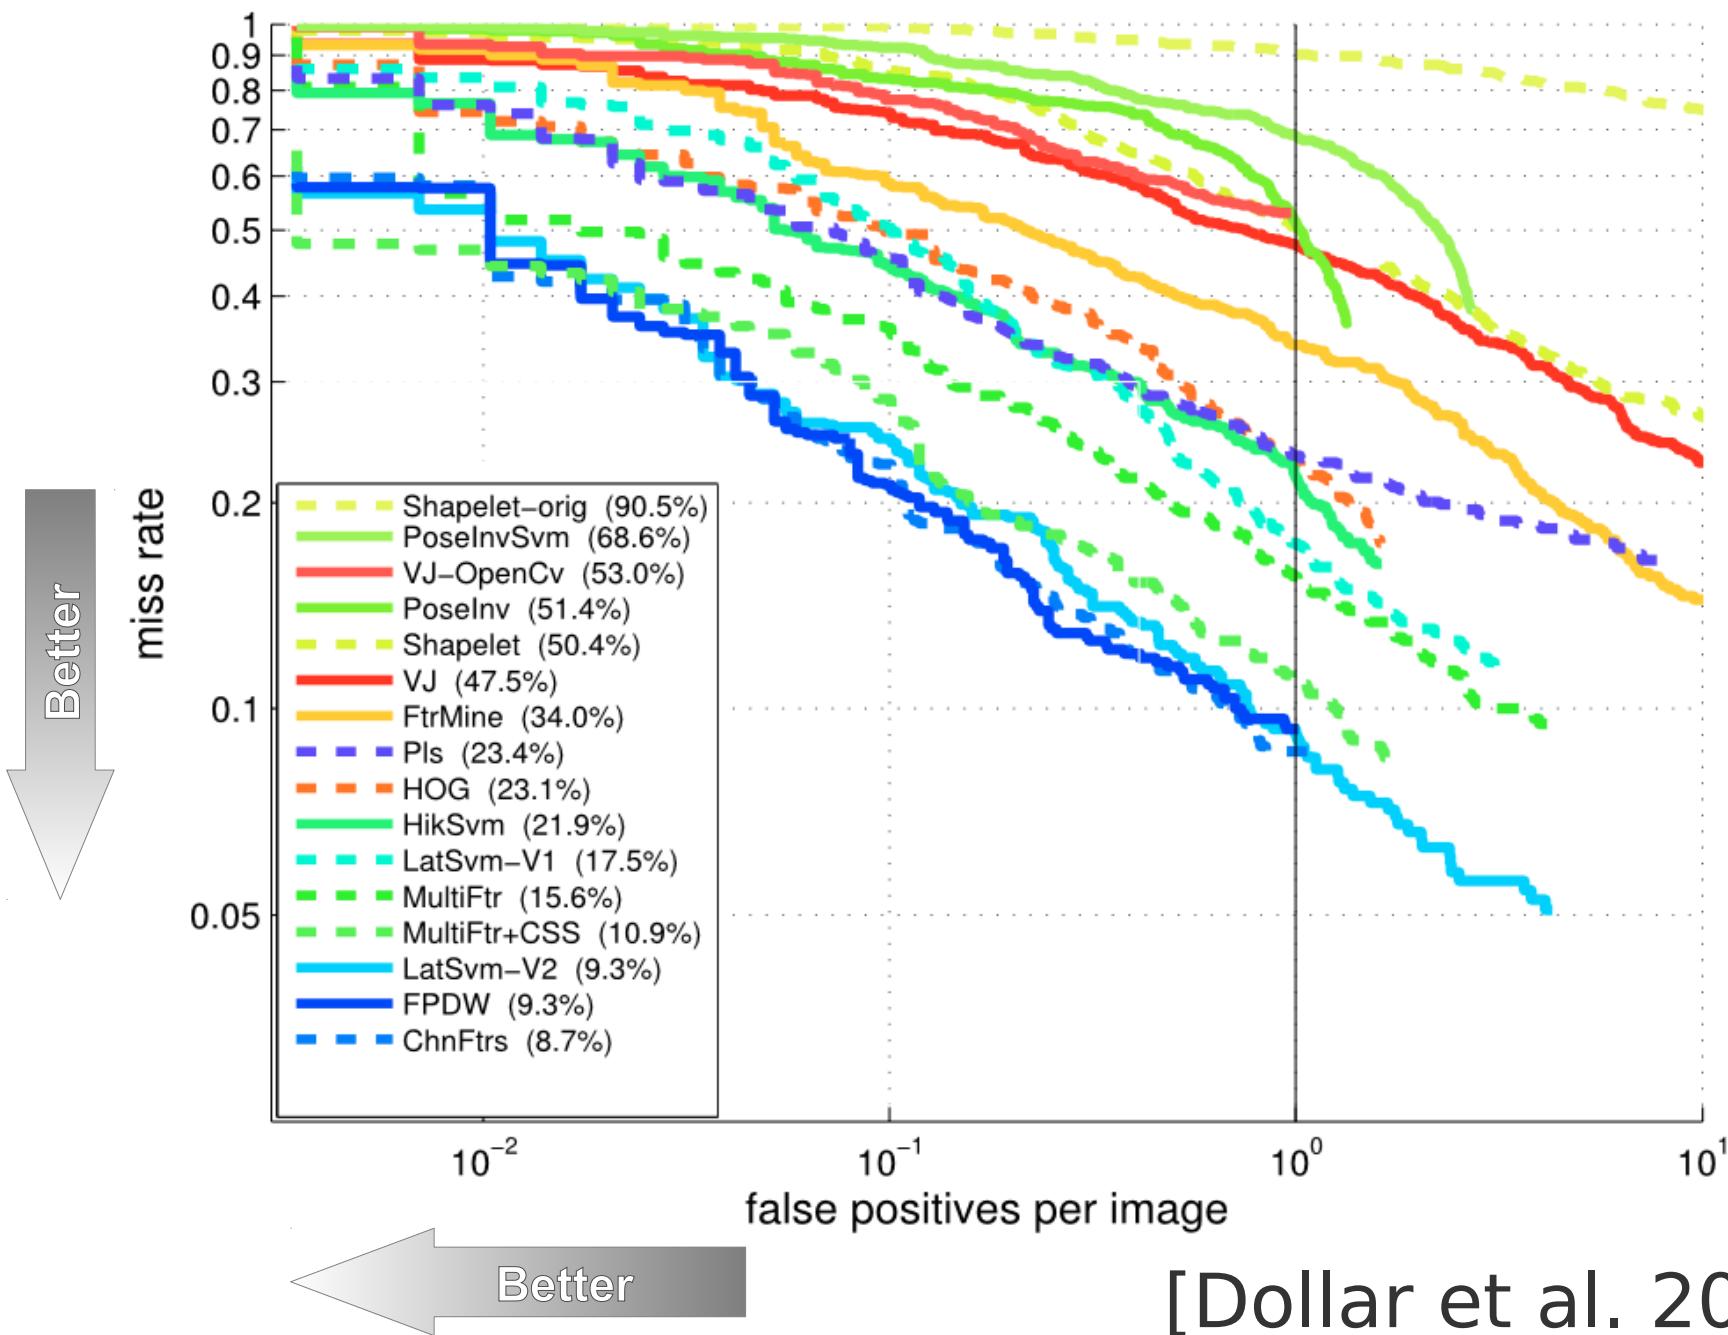


# Why 165 Hz ?

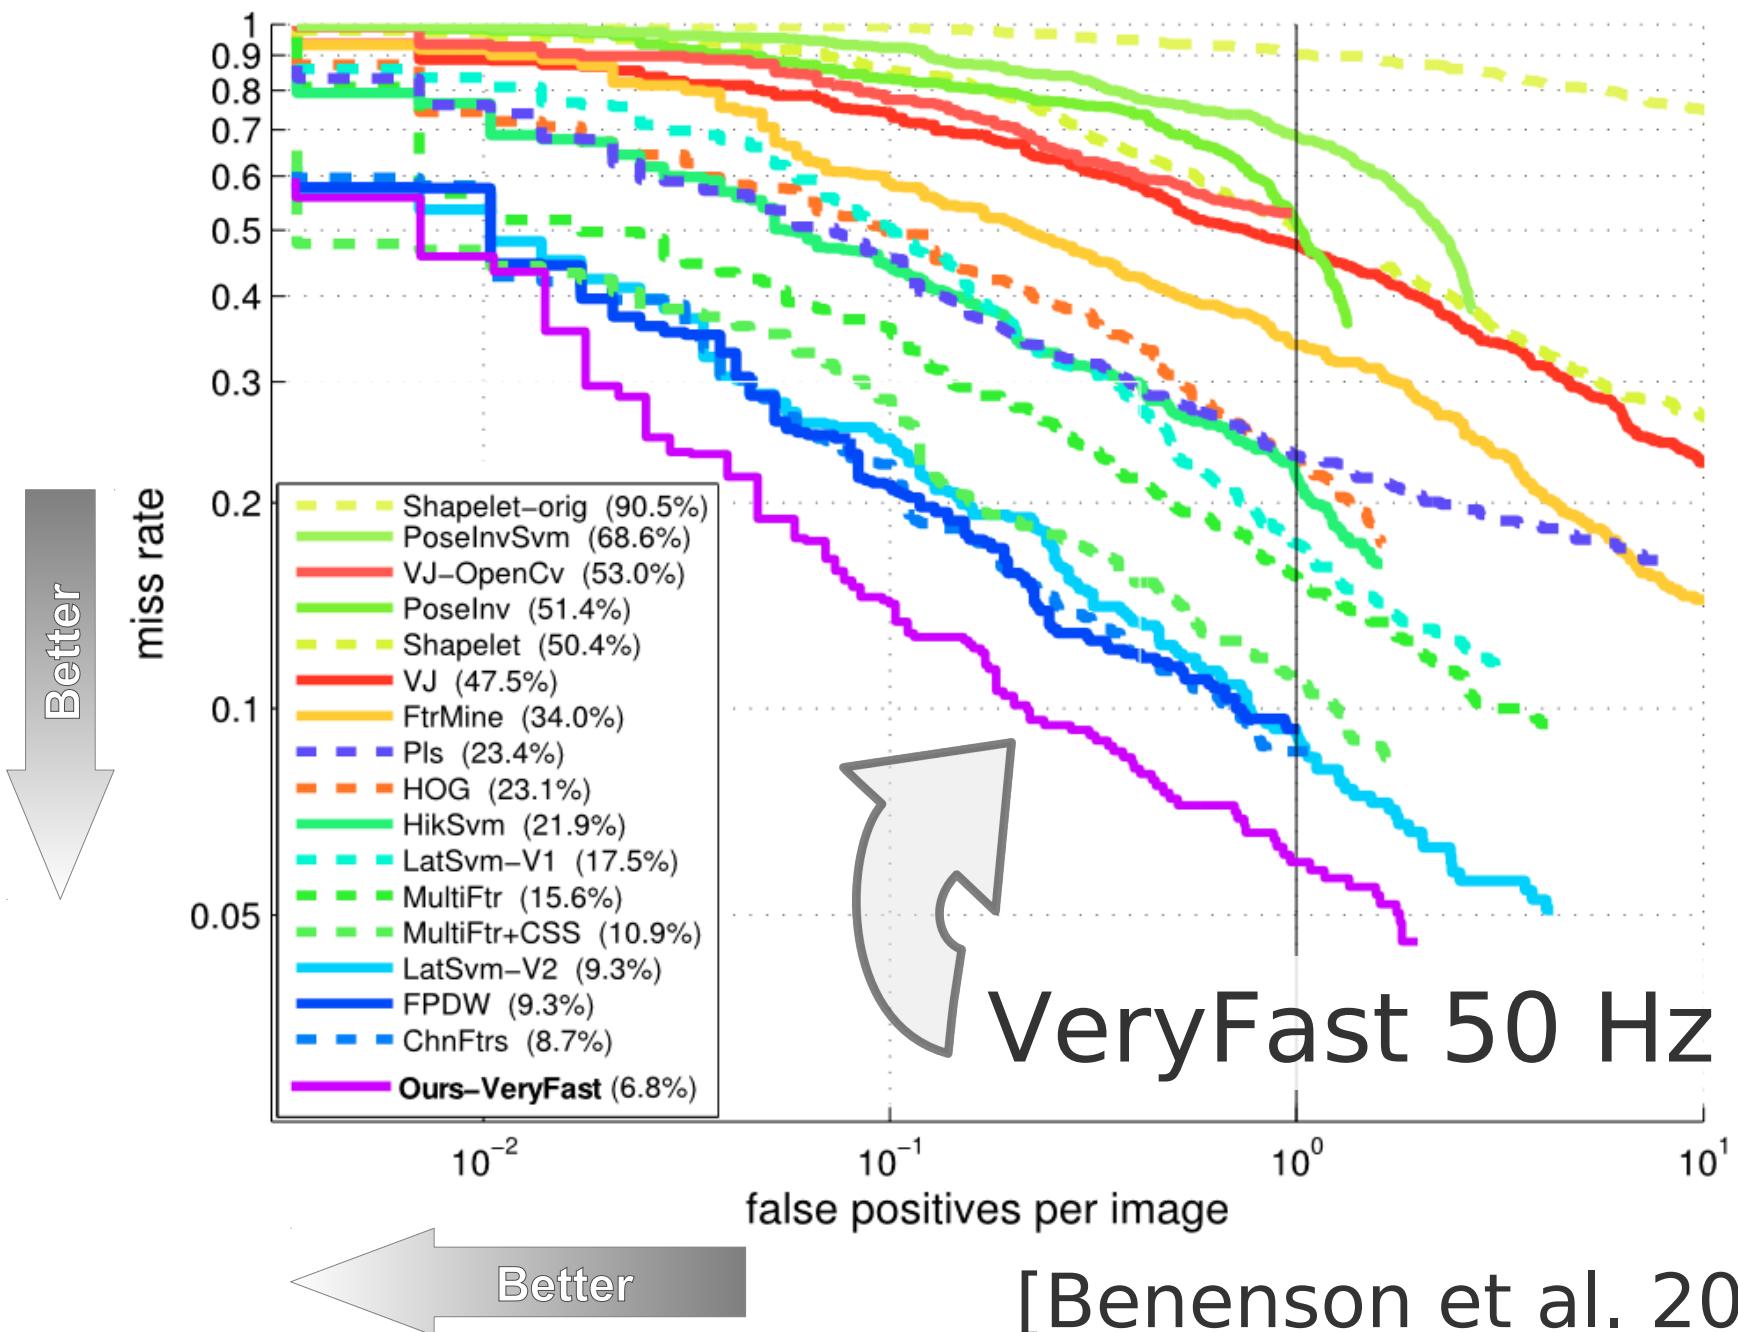
- Detection is one module amongst many
- Less computational power
- *Latency matters*



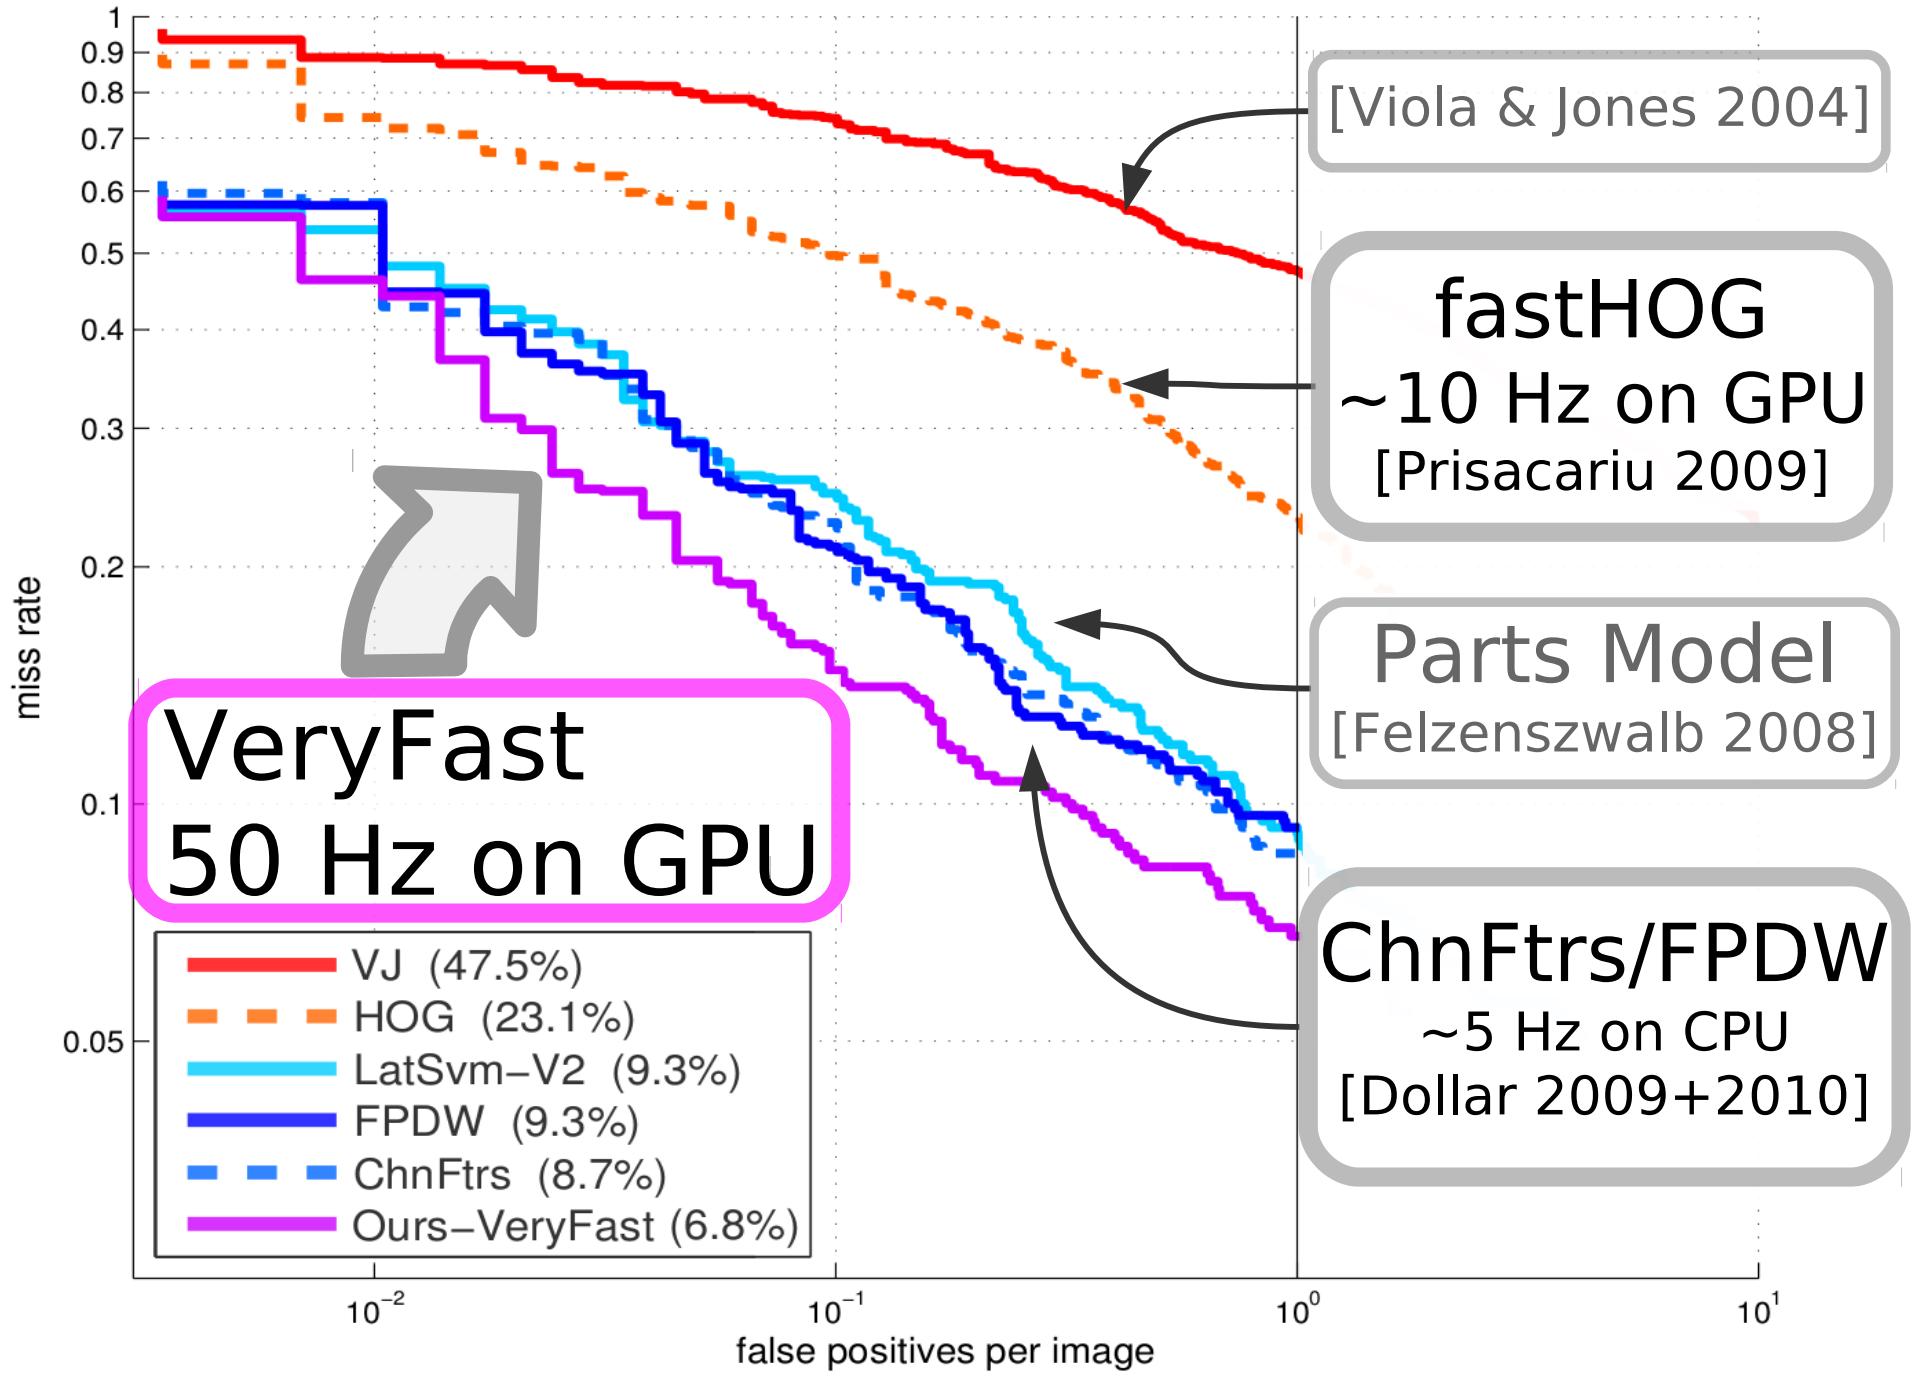
# INRIA dataset



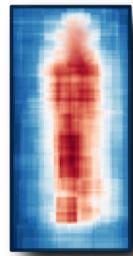
# INRIA dataset



INRIA dataset



# Detecting objects without image resizing



1 model,  
50 image scales

50 models,  
1 image scale

[Benenson et al. 2012]

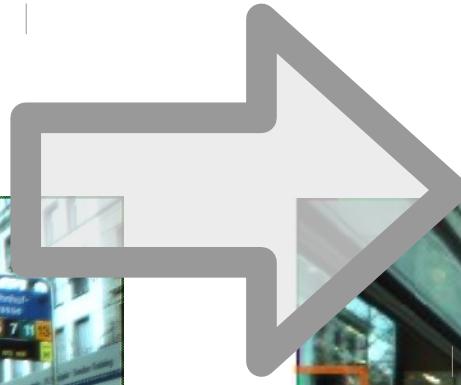
# We want to use scene geometry to guide the detections



Monocular

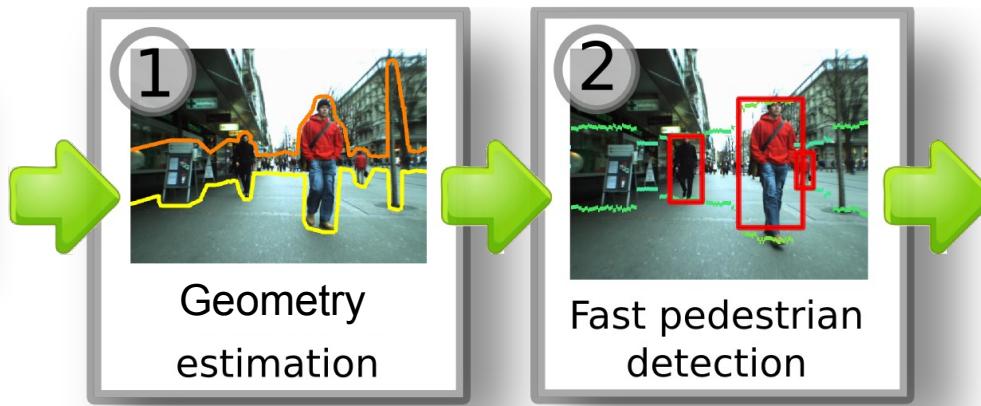


Stereo





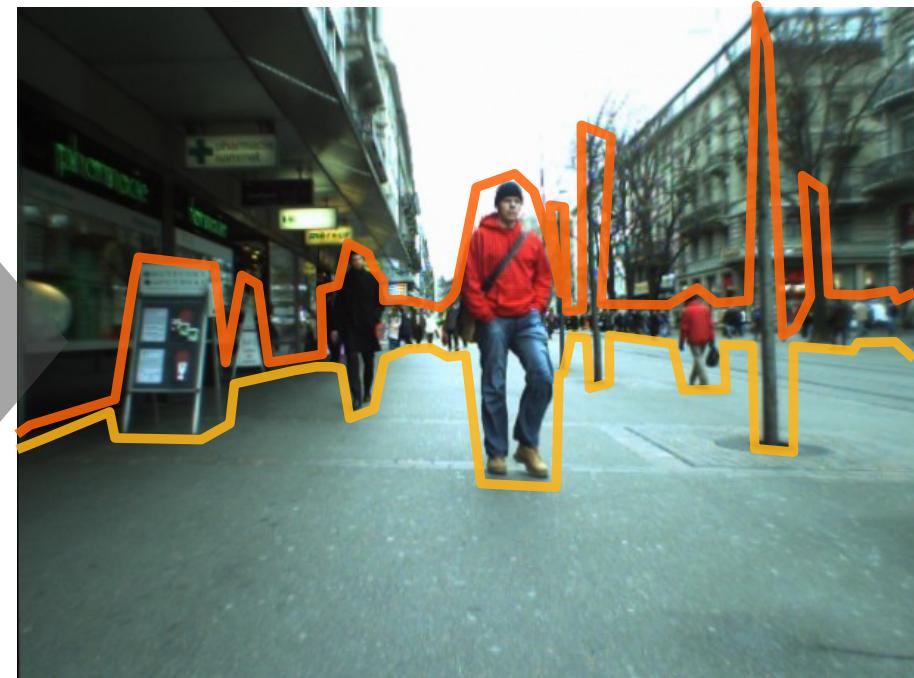
Input stereo image



Detected  
pedestrians

>50 Hz on GPU

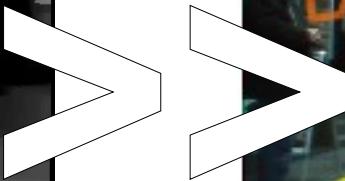
# Depth maps are slow to compute



<50 Hz on CPU

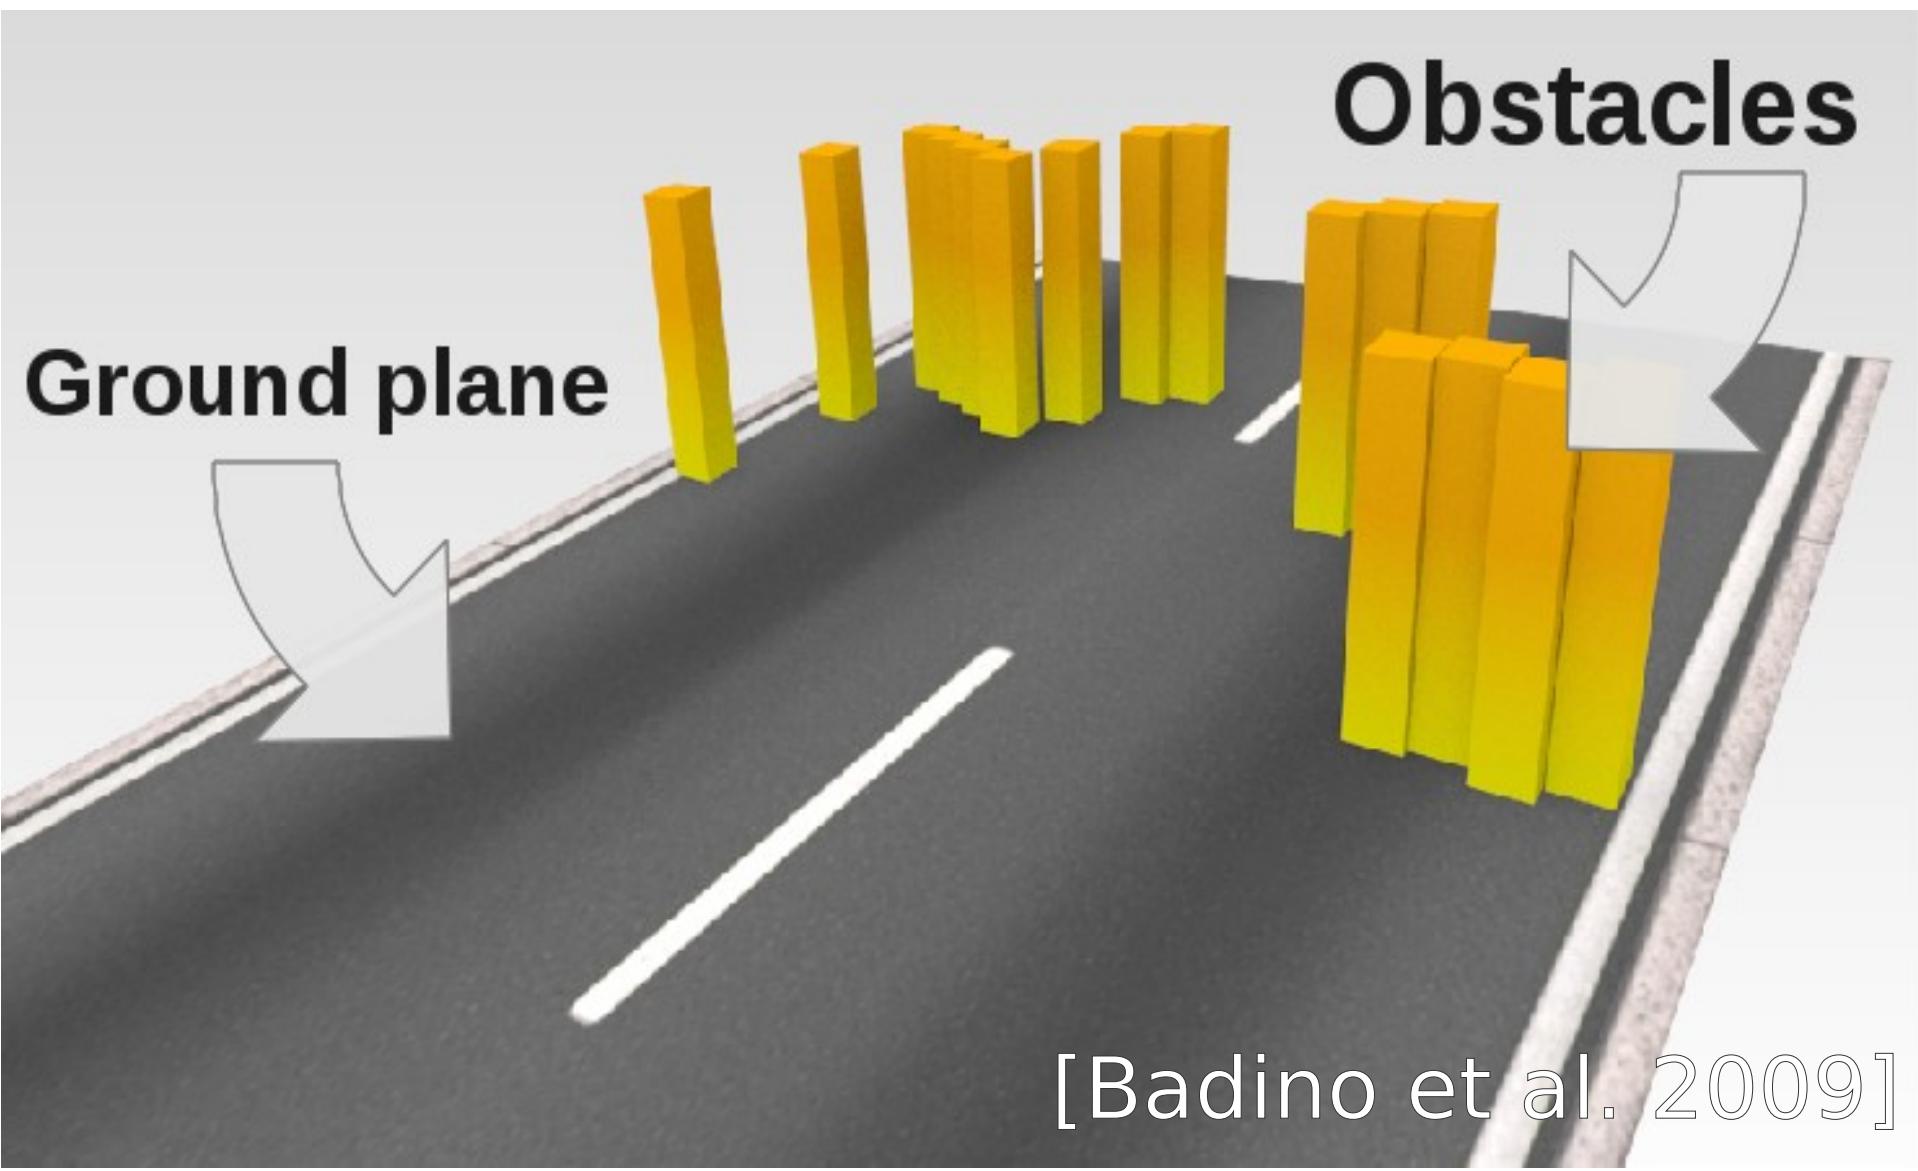
What to do when  
depth map computation  
is much slower than detection?

# Depth maps compute more information than we need



Key insight

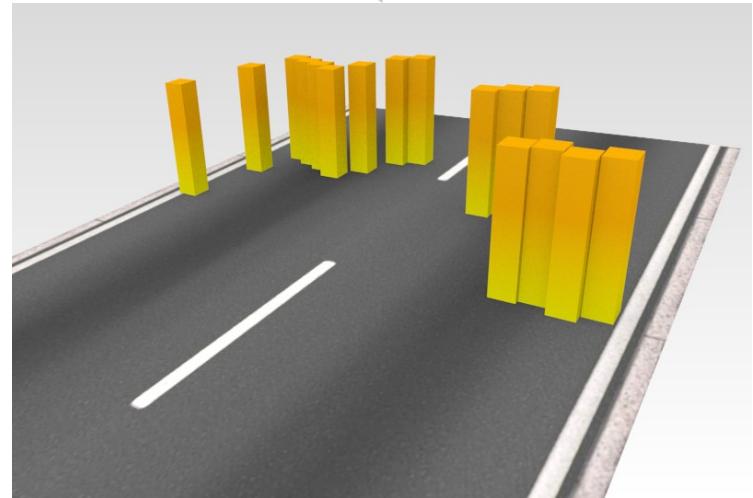
# Stixel world



# Stixel world without depth map



90 Hz on CPU

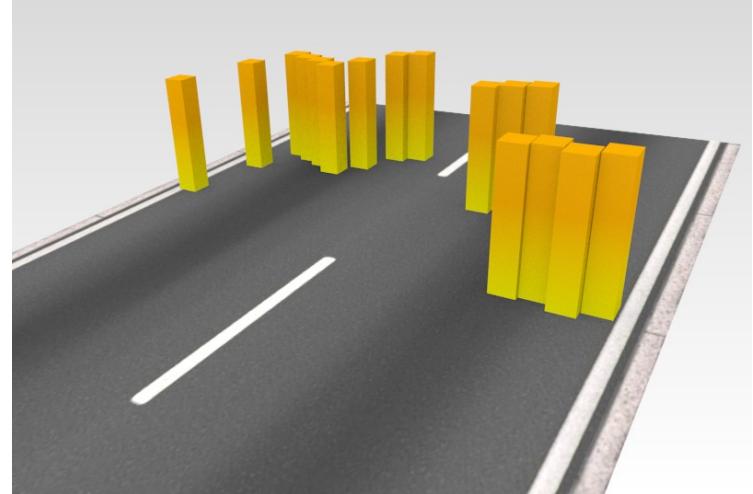


[Benenson et al. 2011]

# Stixel world without depth map



135 Hz  
→ 90 Hz on CPU

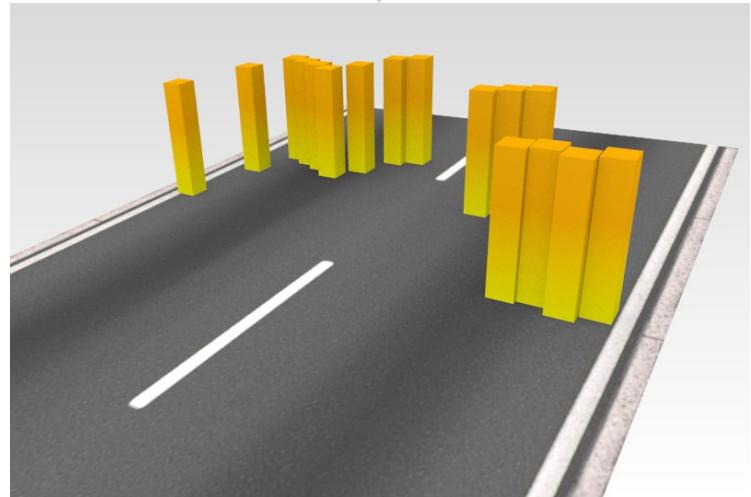


[Benenson et al. 2011 & 2012]

# Stixel world without depth map



**300 Hz on CPU**



## CVVT@ICCV 2011

“we can extract scene geometry from stereo at 90 Hz”

## CVPR 2012

“we can detect pedestrian at 50 Hz,  
135 Hz when using scene geometry”

## CVVT@ECCV 2012

“we can extract **scene geometry from stereo at ~300 Hz**,  
and thus detect pedestrian at 165 Hz”

Is (direct) stixel world a good idea?

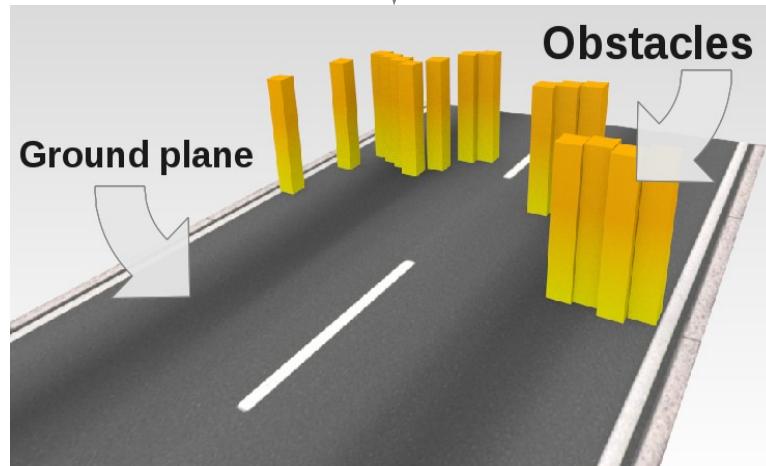
# It opens new possibilities



Monocular

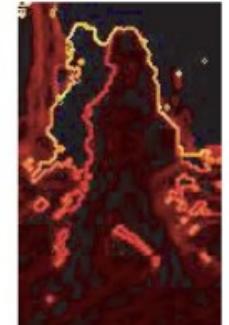
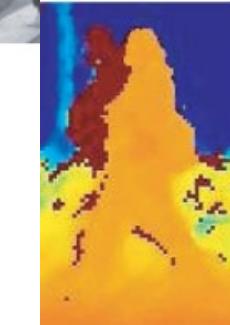
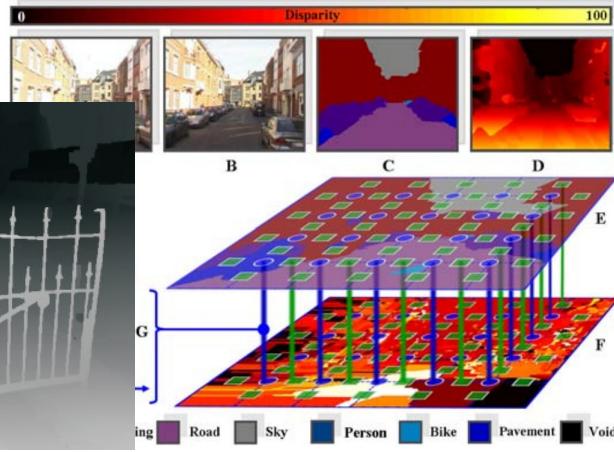


Pixel-wise  
depth map



Stixel world

# Stereo and object detection is a chicken and egg problem!



[Ladicky et al. 2010], [Bleyer et al. 2011]

# Stixel world without depth map

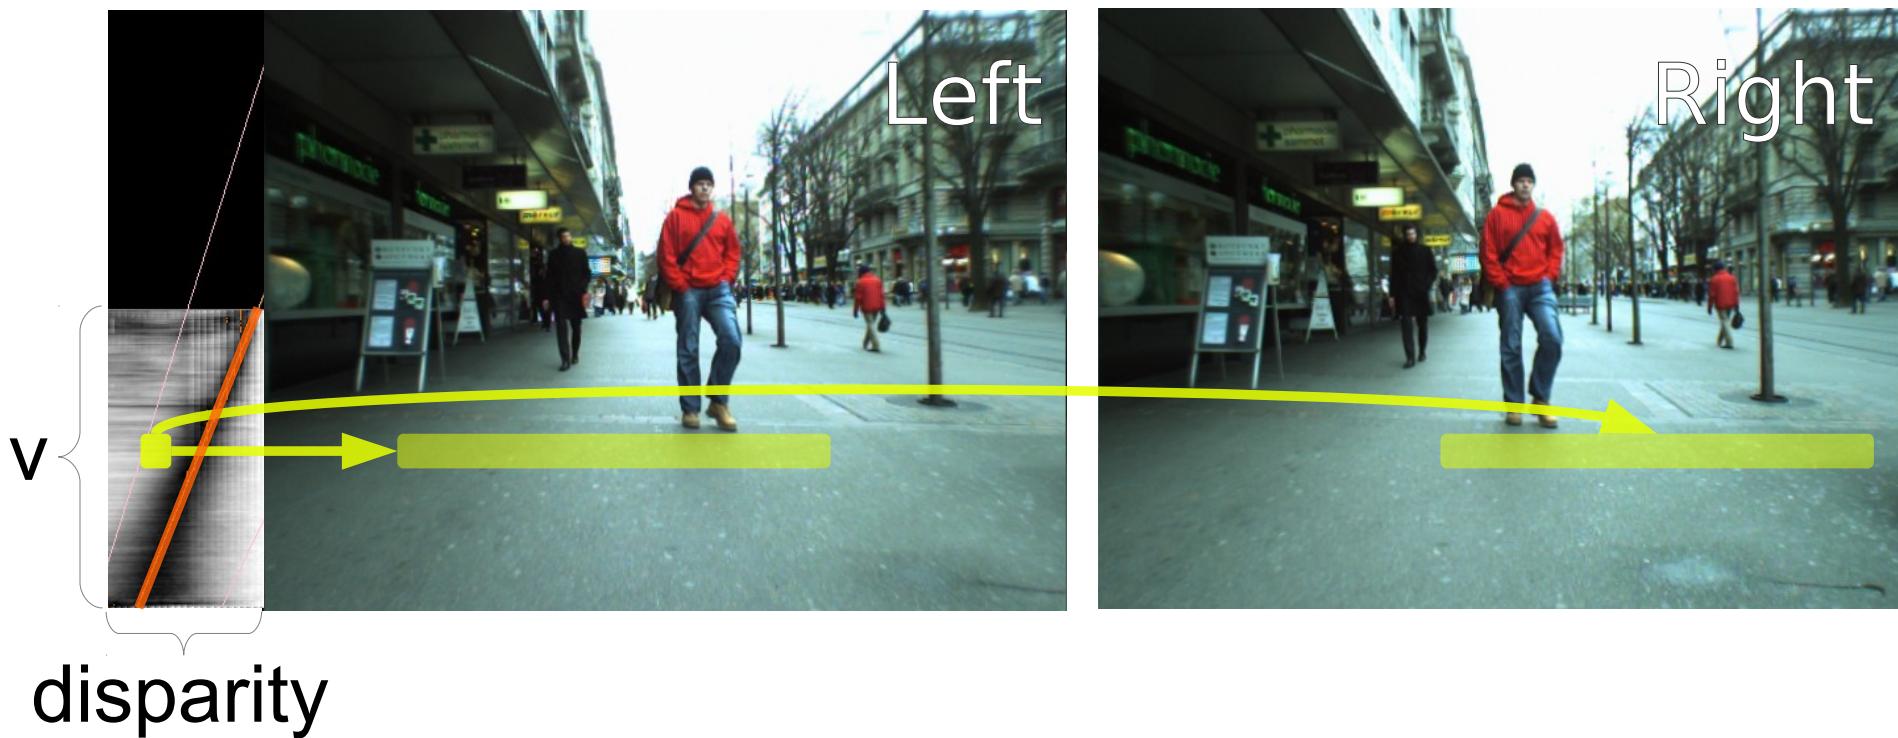


# Stixel world without depth map



1) Ground plane estimation

# Stixel world without depth map



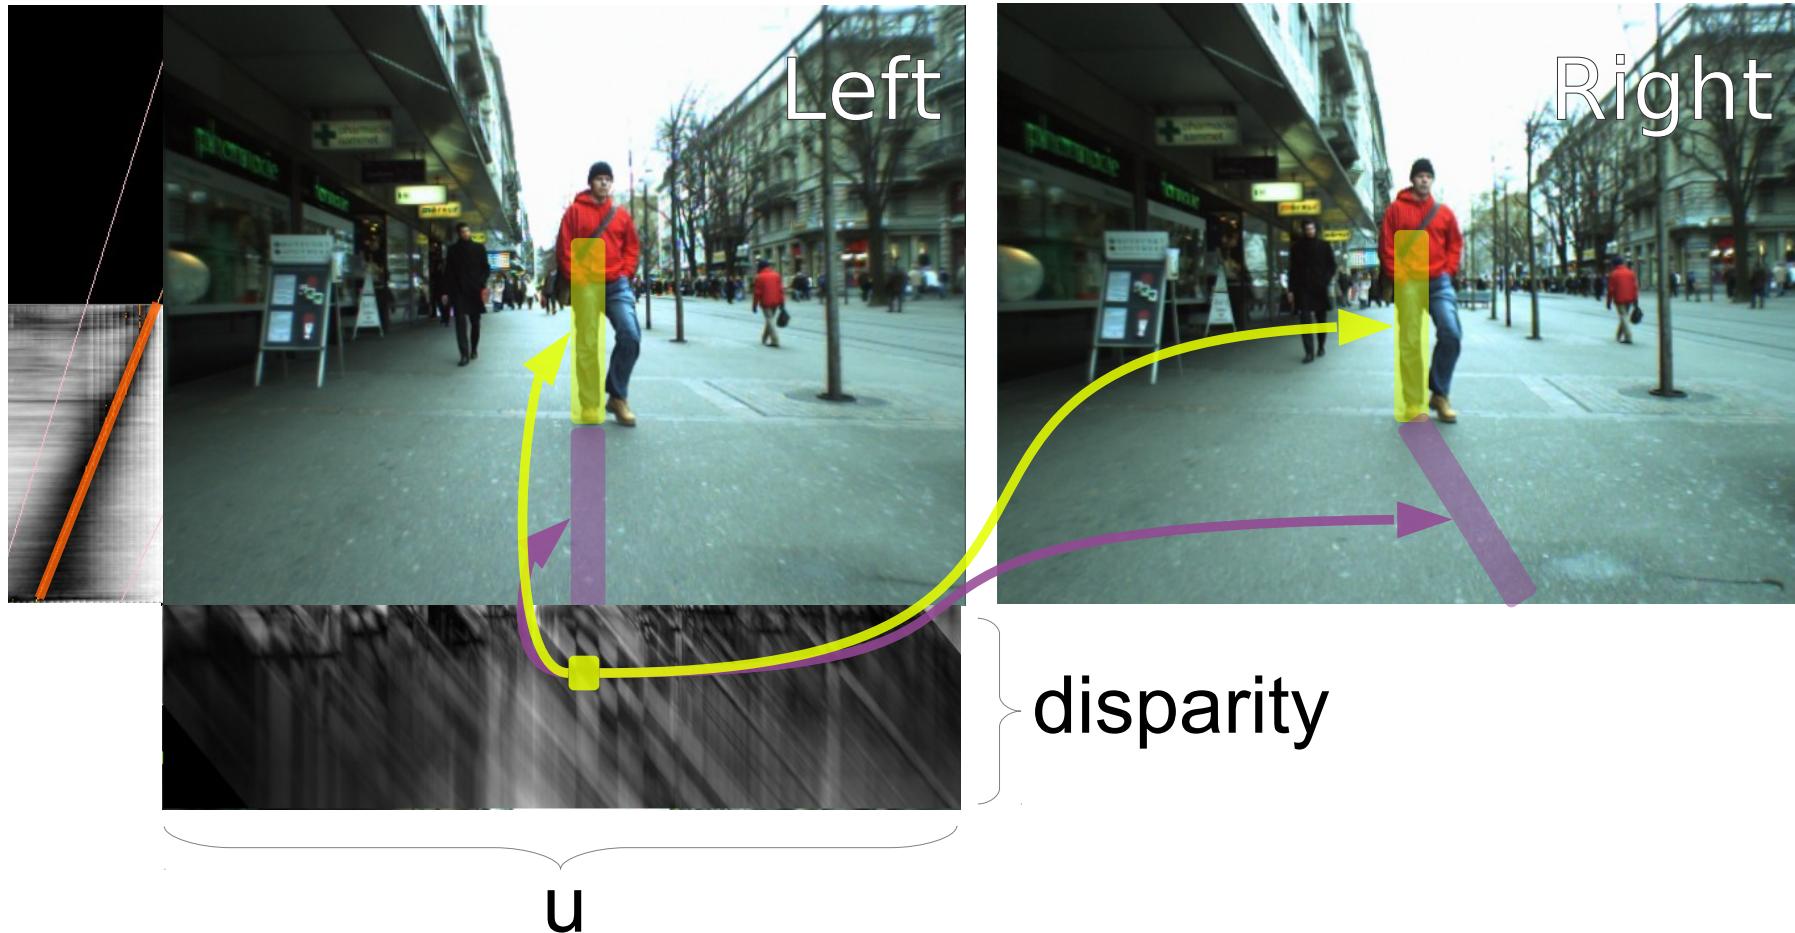
1) Ground plane estimation

# Stixel world without depth map



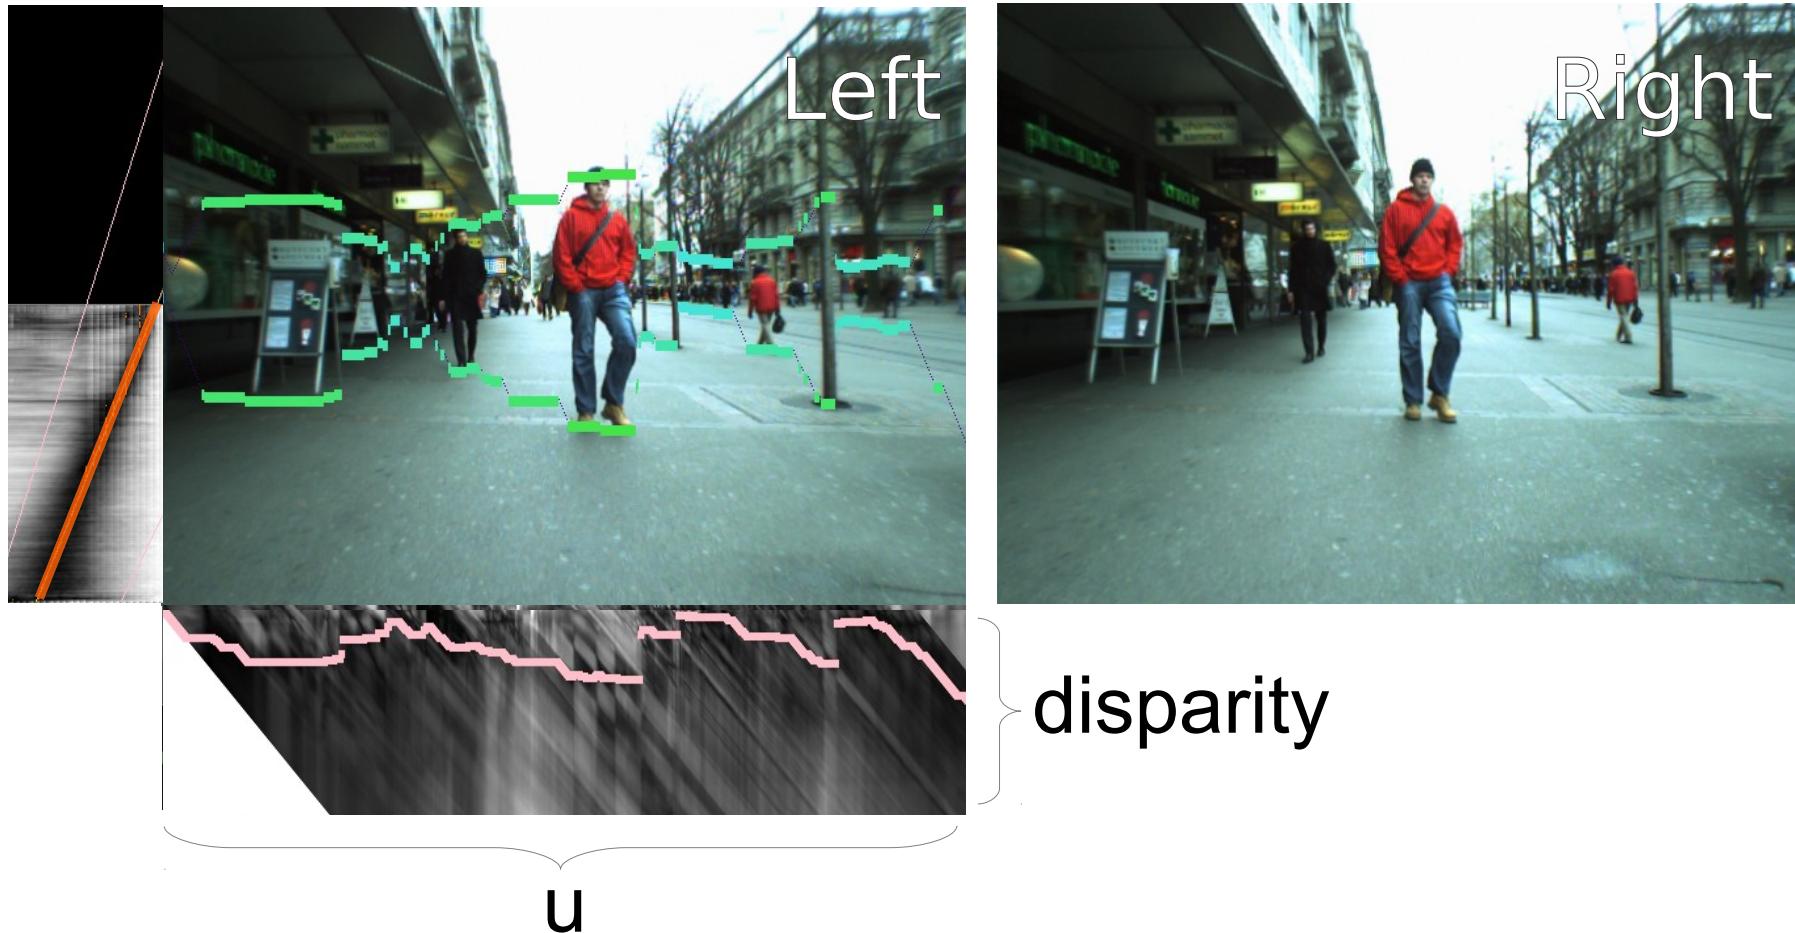
2) Stixel distance estimation

# Stixel world without depth map



2) Stixel distance estimation

# Stixel world without depth map



2) Stixel distance estimation @ 135 Hz CPU

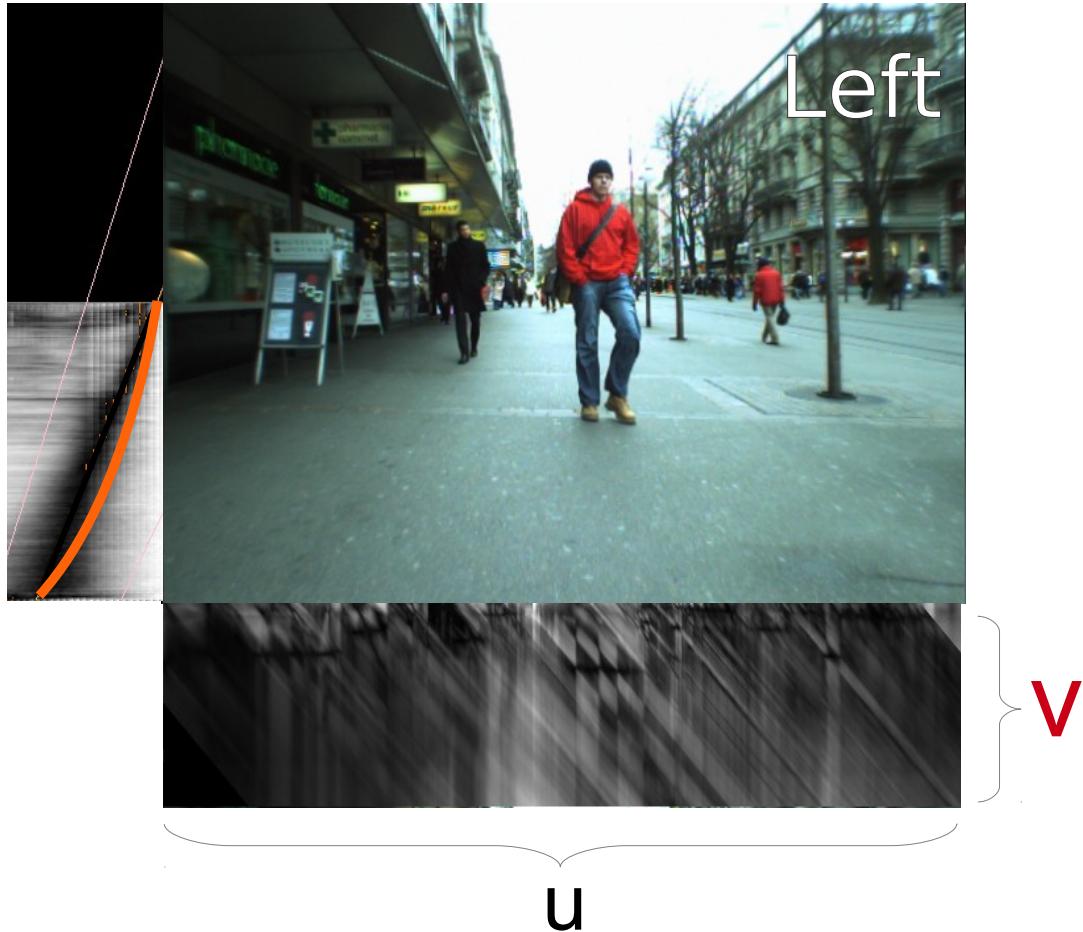
**100 Hz is too slow!**

Our initial proposal has  
(at least) two issues

# Too focused in u-disparity, instead of the image space

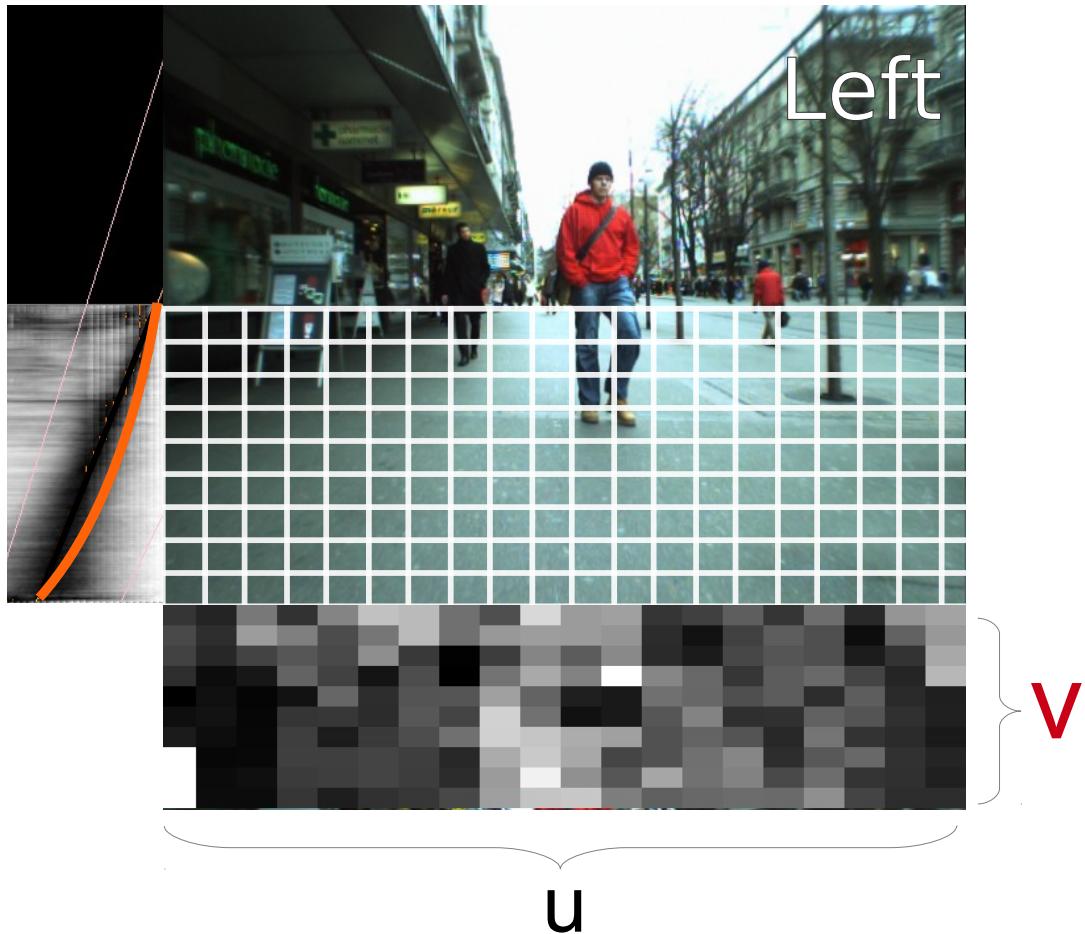


# Too focused in u-disparity, instead of the image space



We will **quantize** both rows and columns

# Too focused in u-disparity, instead of the u-v space

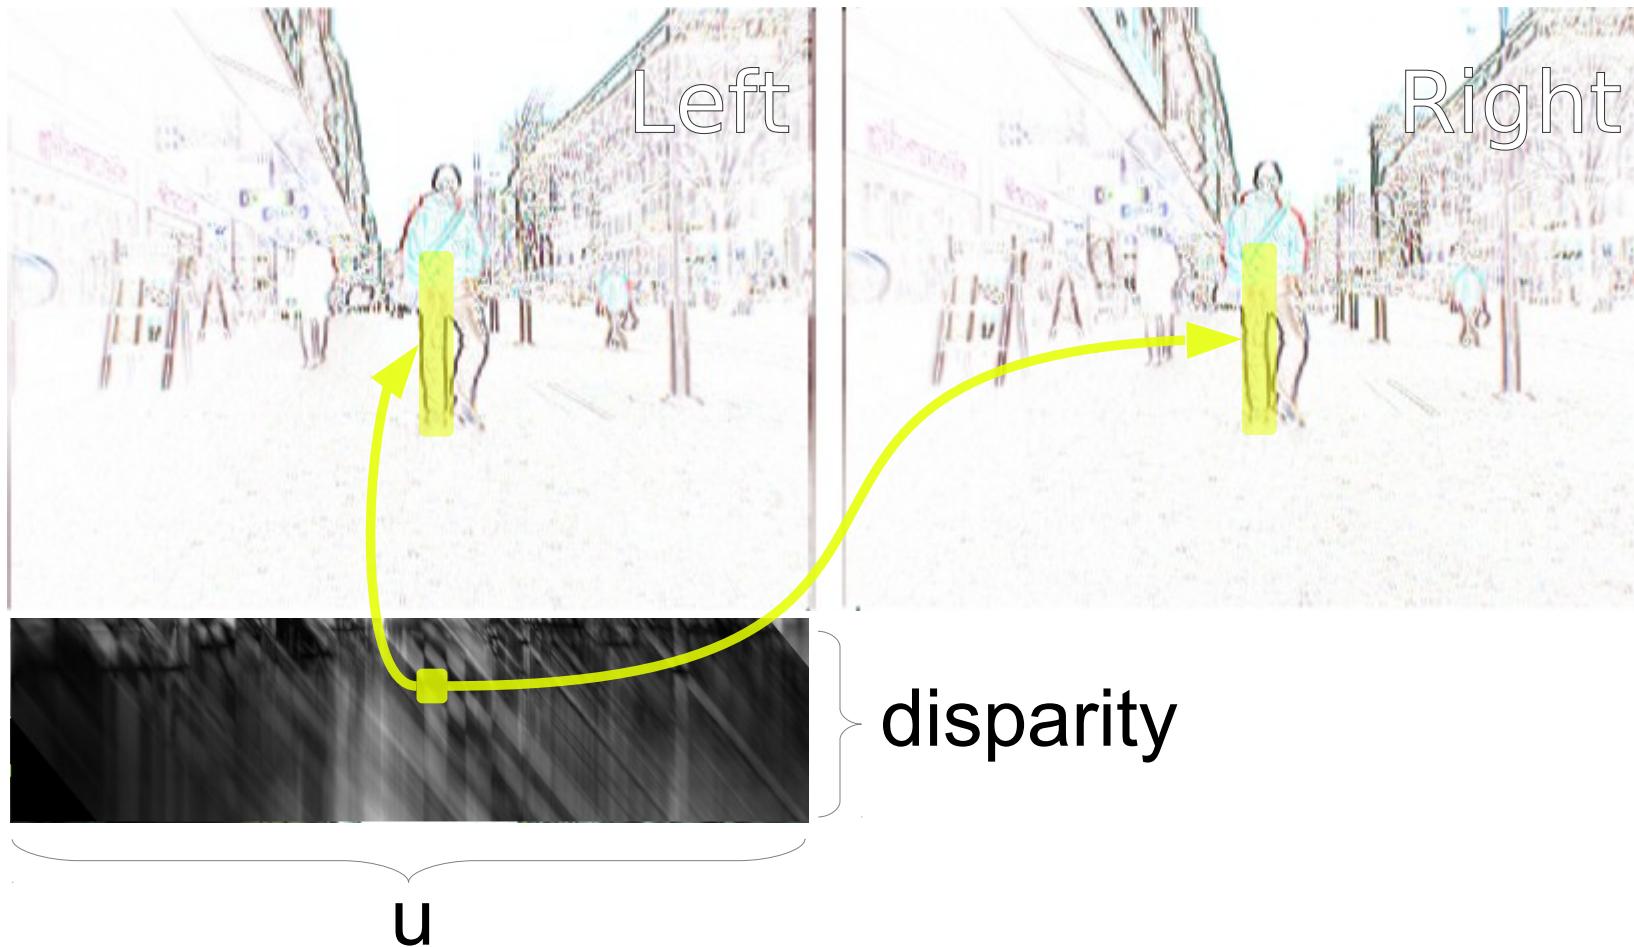


We will **quantize** both rows and columns

# Only half of the image information is used



# Only half of the image information is used



Only gradients along rows are used

# Only half of the image information is used

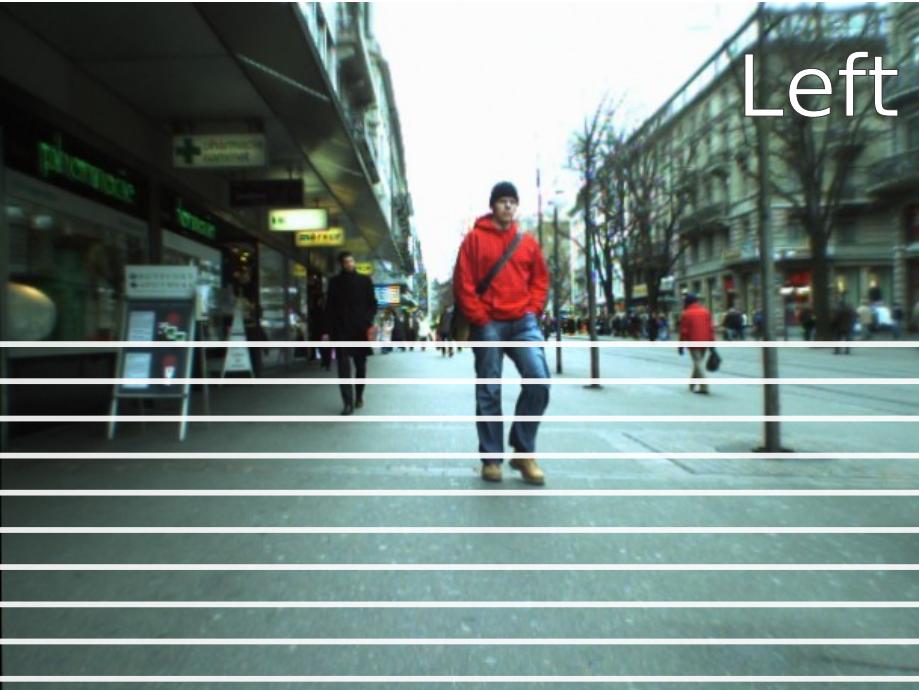


Gradients along columns are **ignored**

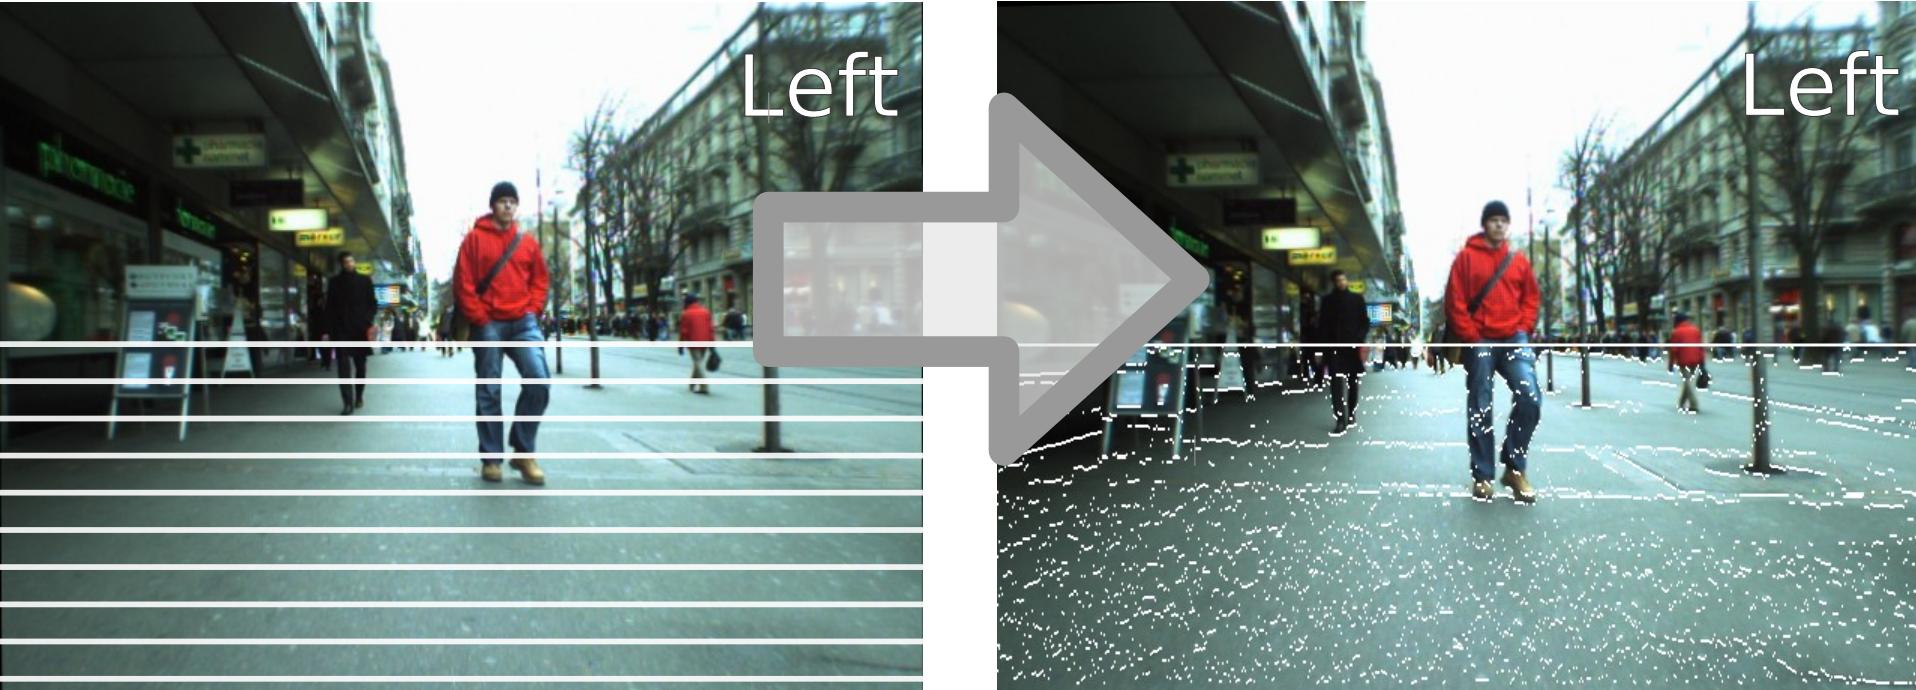
# Now we fully exploit the left image



# Now we fully exploit the left image

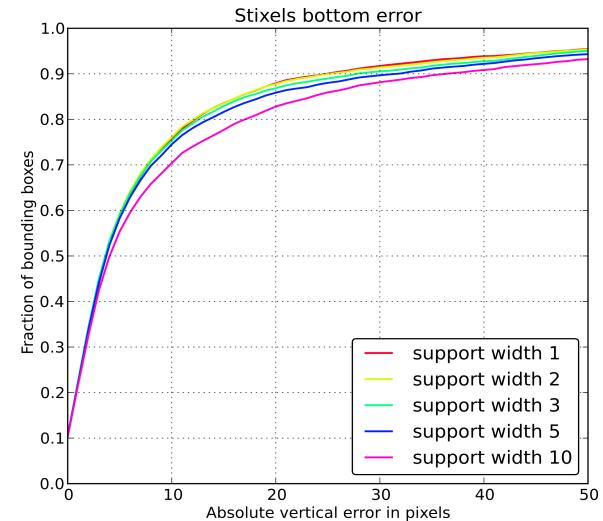
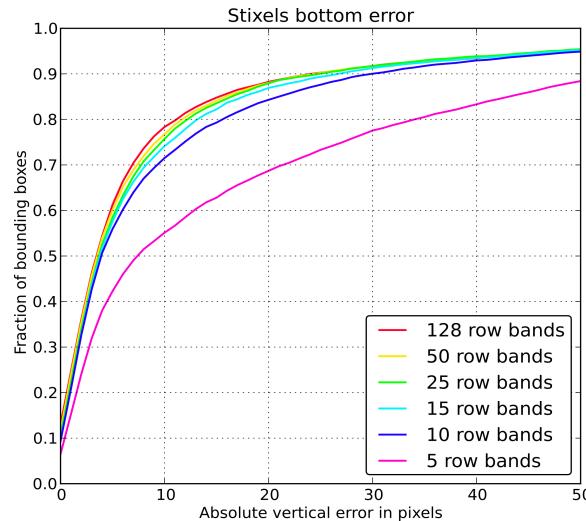
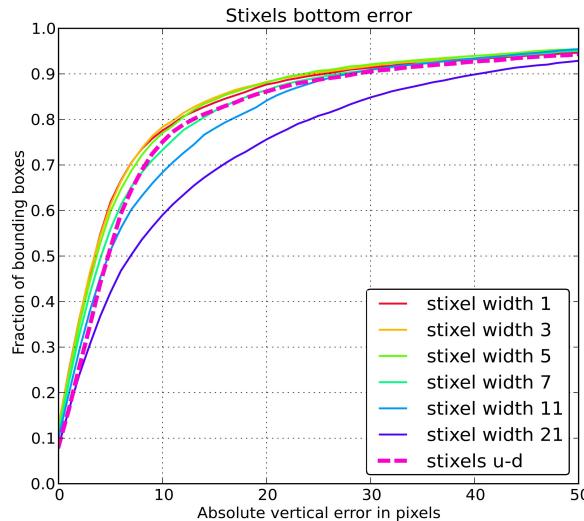


# Now we fully exploit the left image



# What happens when .... ?

- Stixels vertical sampling changes?
- Stixels horizontal sampling changes?
- Stixel support width changes?

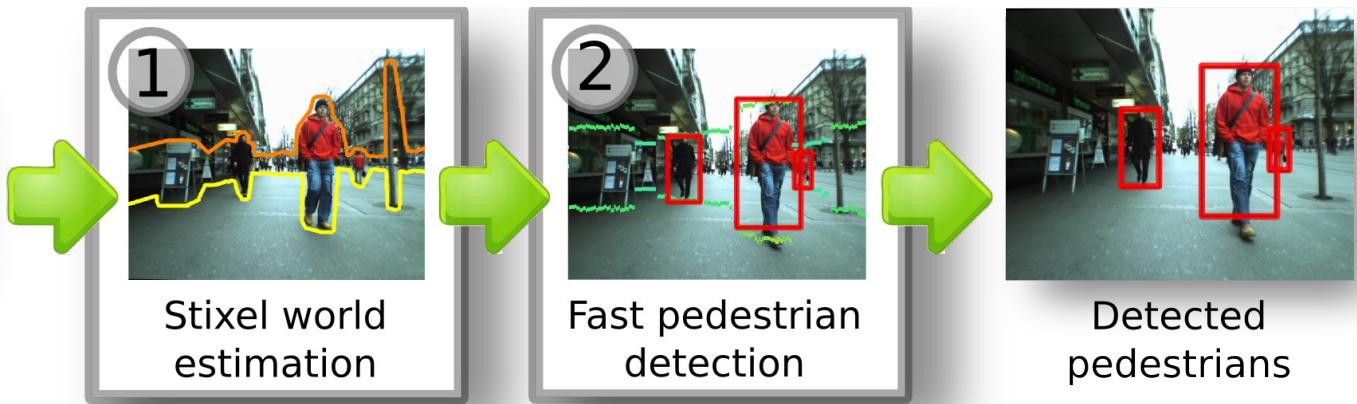


# What happens when .... ?

- Stixels vertical sampling changes?  
→ each **3 columns** is good enough
  - Stixels horizontal sampling changes?  
→ **25 row bands** is good enough
  - Stixel support width changes?  
→ **width 1 column** is best
- ~100 Hz ⇒ **~300** Hz on CPU  
(on 640x480 pixels images)

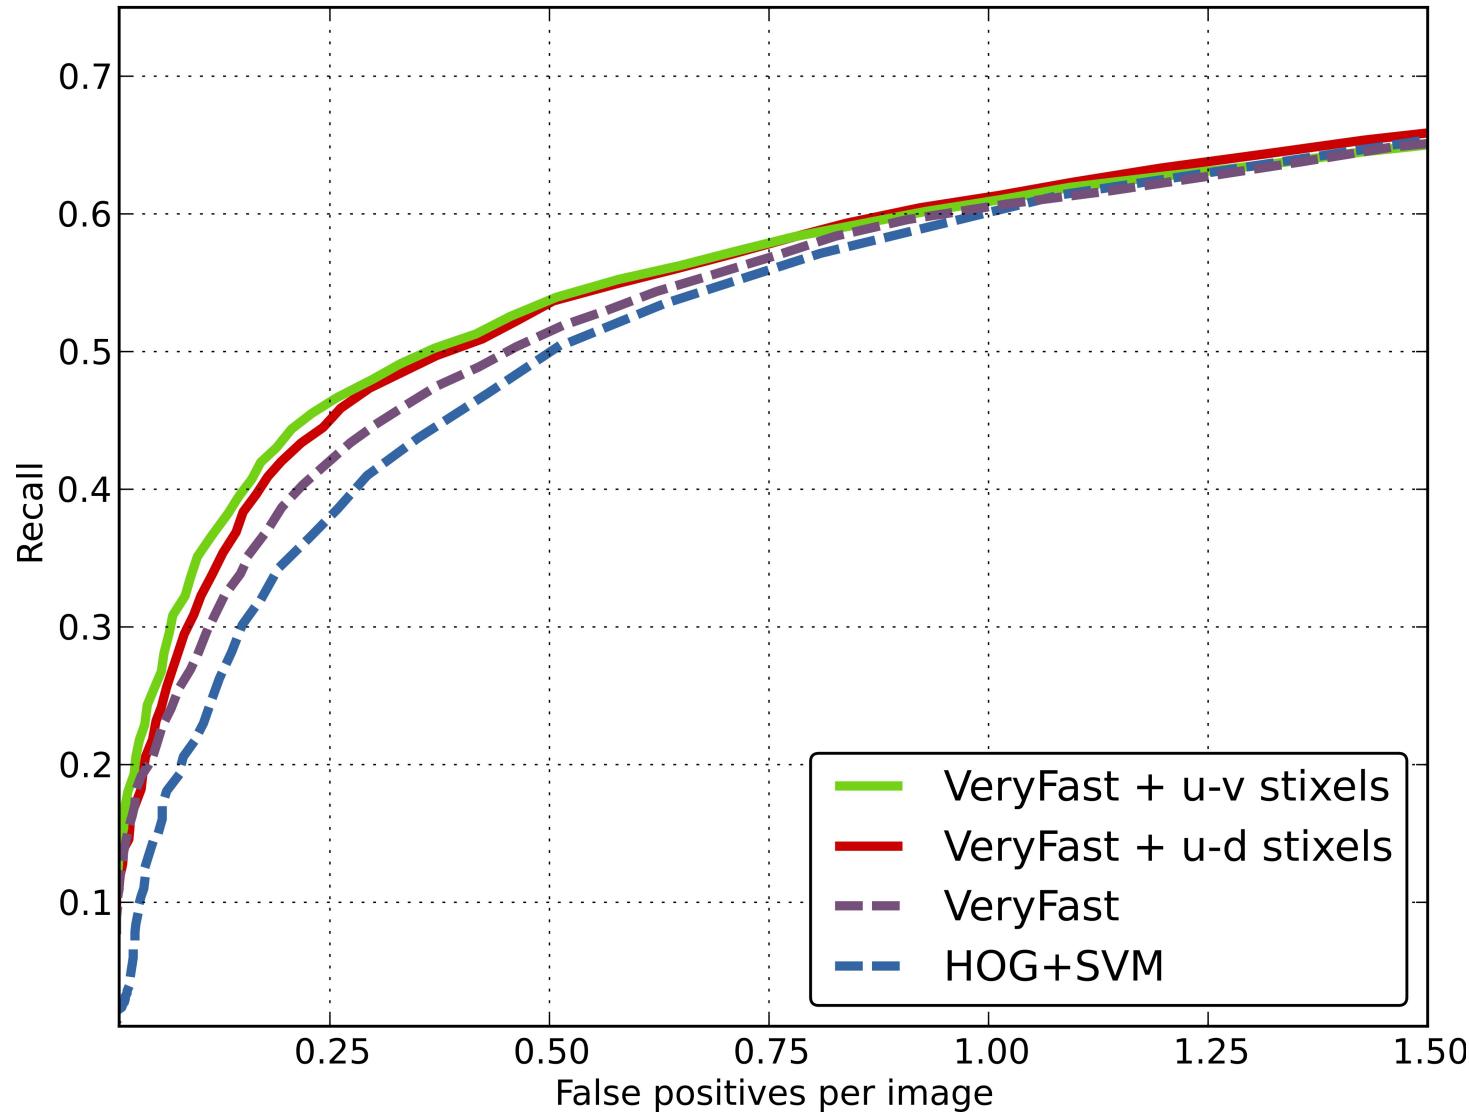


Input stereo image



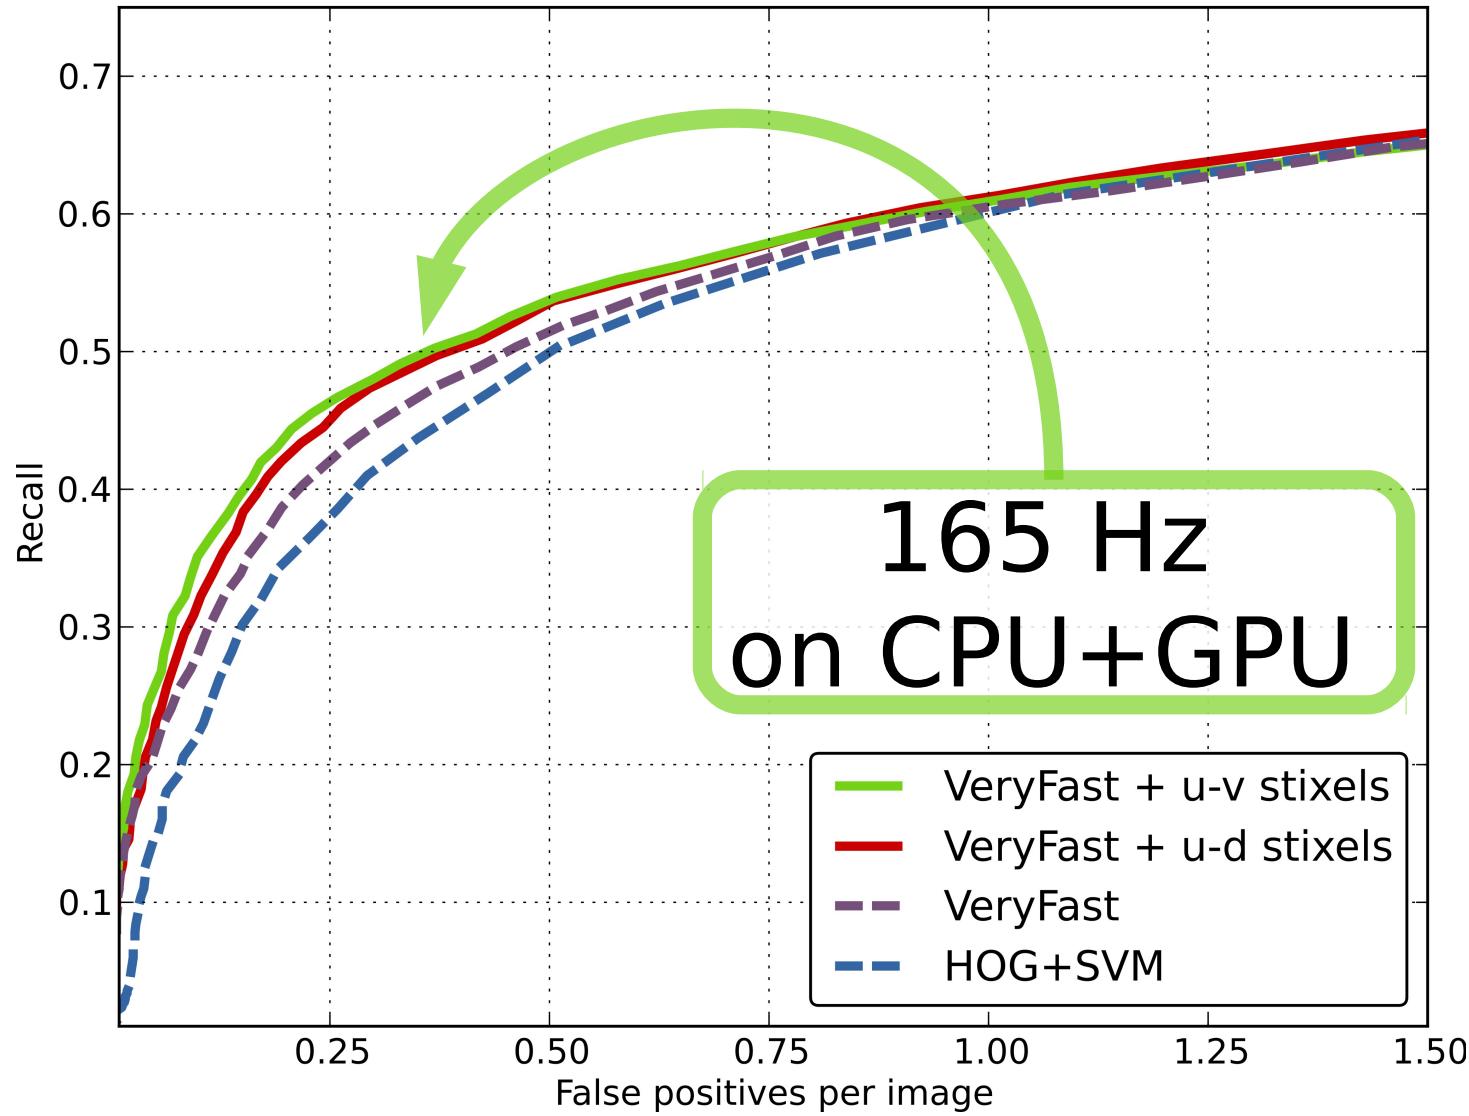
# More speed, same quality

Recall versus FPPI over Bahnhof dataset,  
considering all windows with height > 40 [pixels]

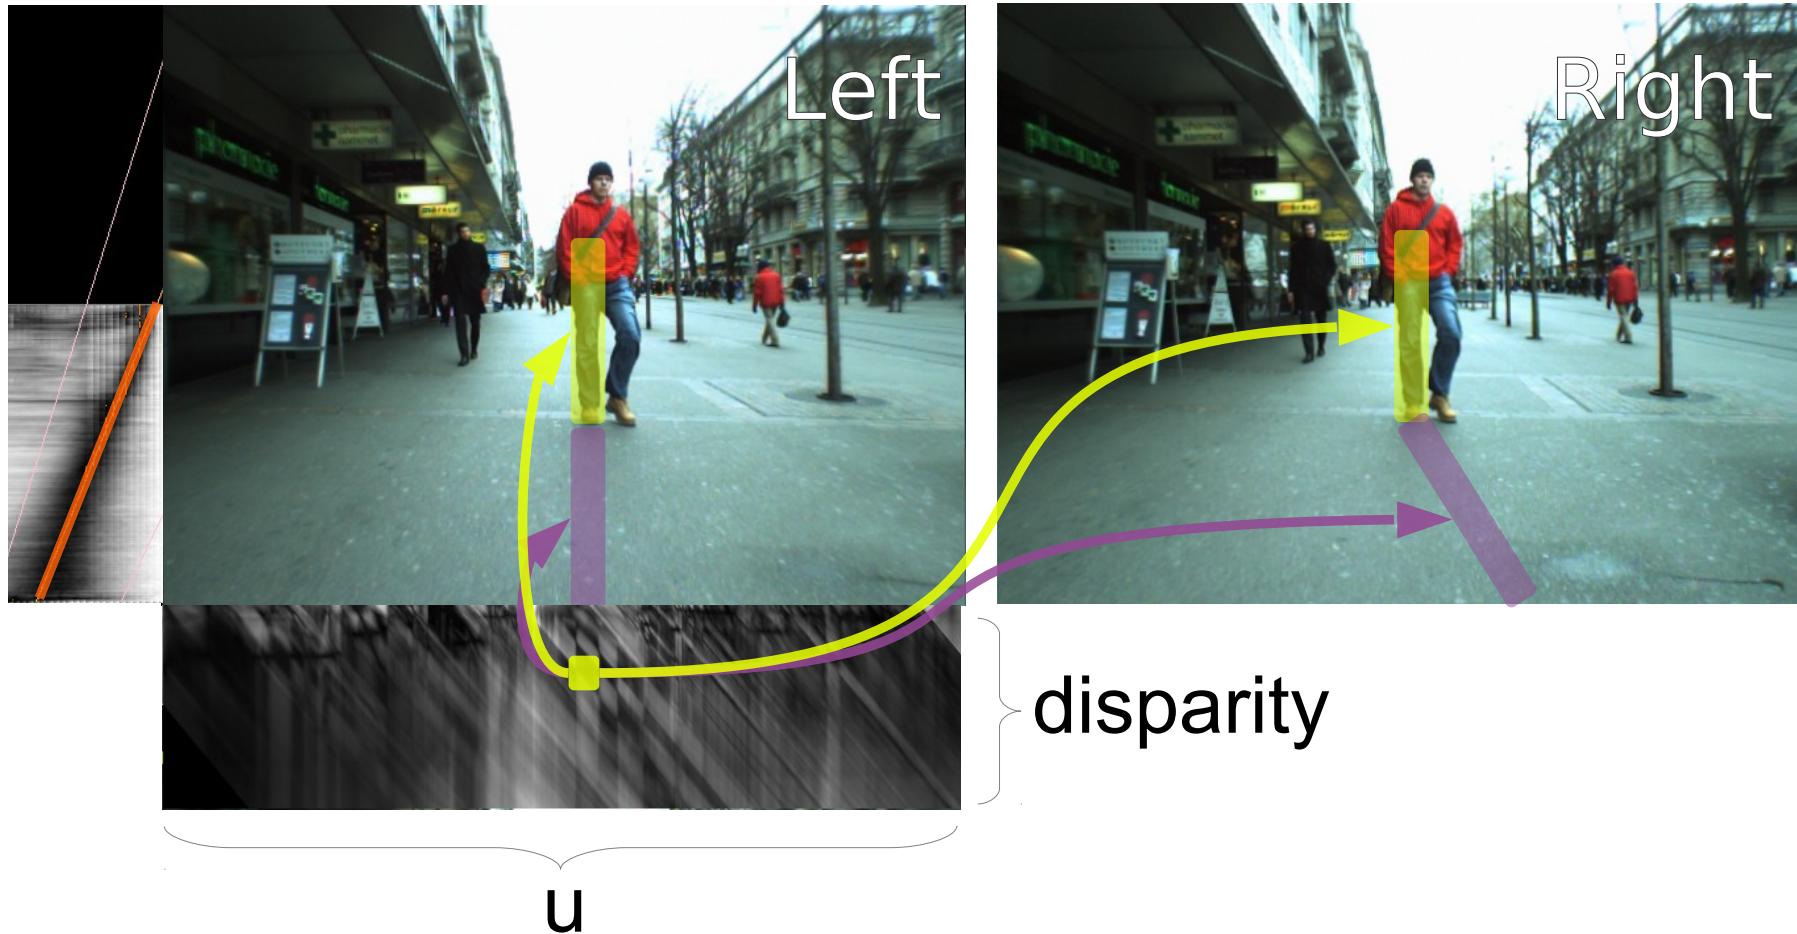


# More speed, same quality

Recall versus FPPI over Bahnhof dataset,  
considering all windows with height > 40 [pixels]

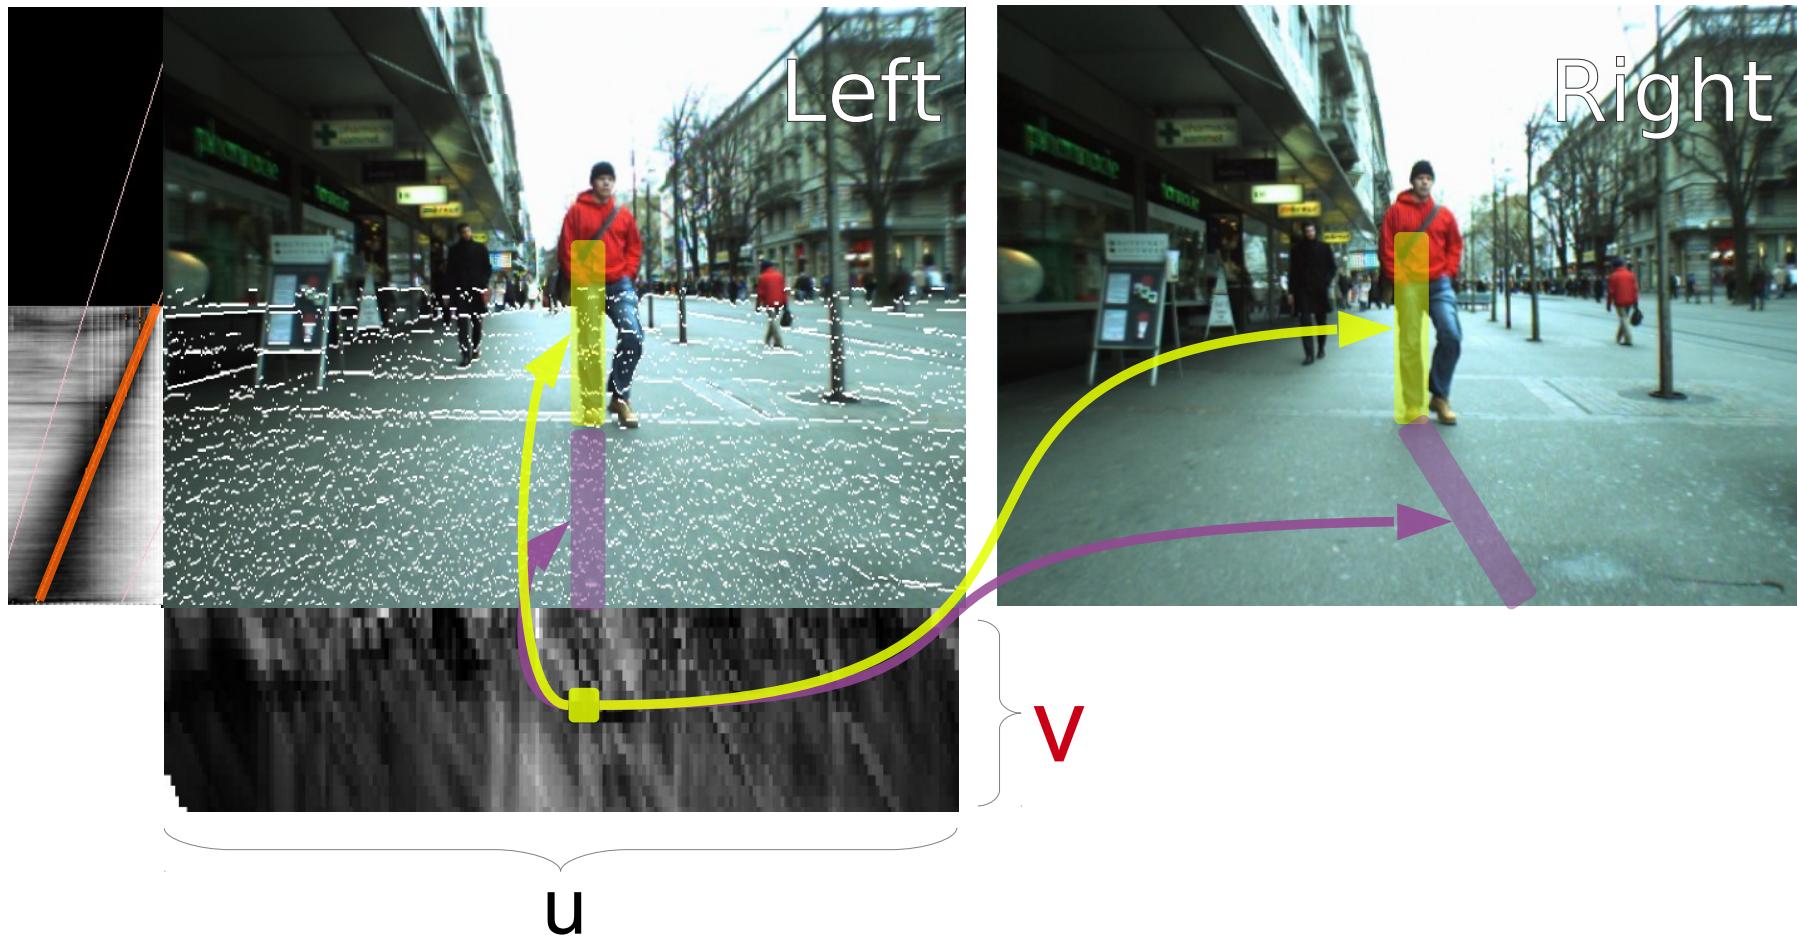


# Stixel estimation in u-disparity space is good



Stixel distance estimation @ 135 Hz CPU

# Stixel estimation in u-v space is better



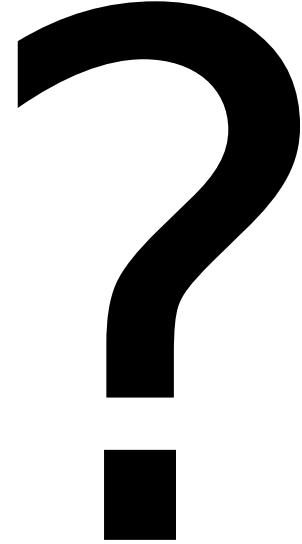
2) Stixel distance estimation **@ 300 Hz CPU**



Take-away message:

Stixels are useful  
and

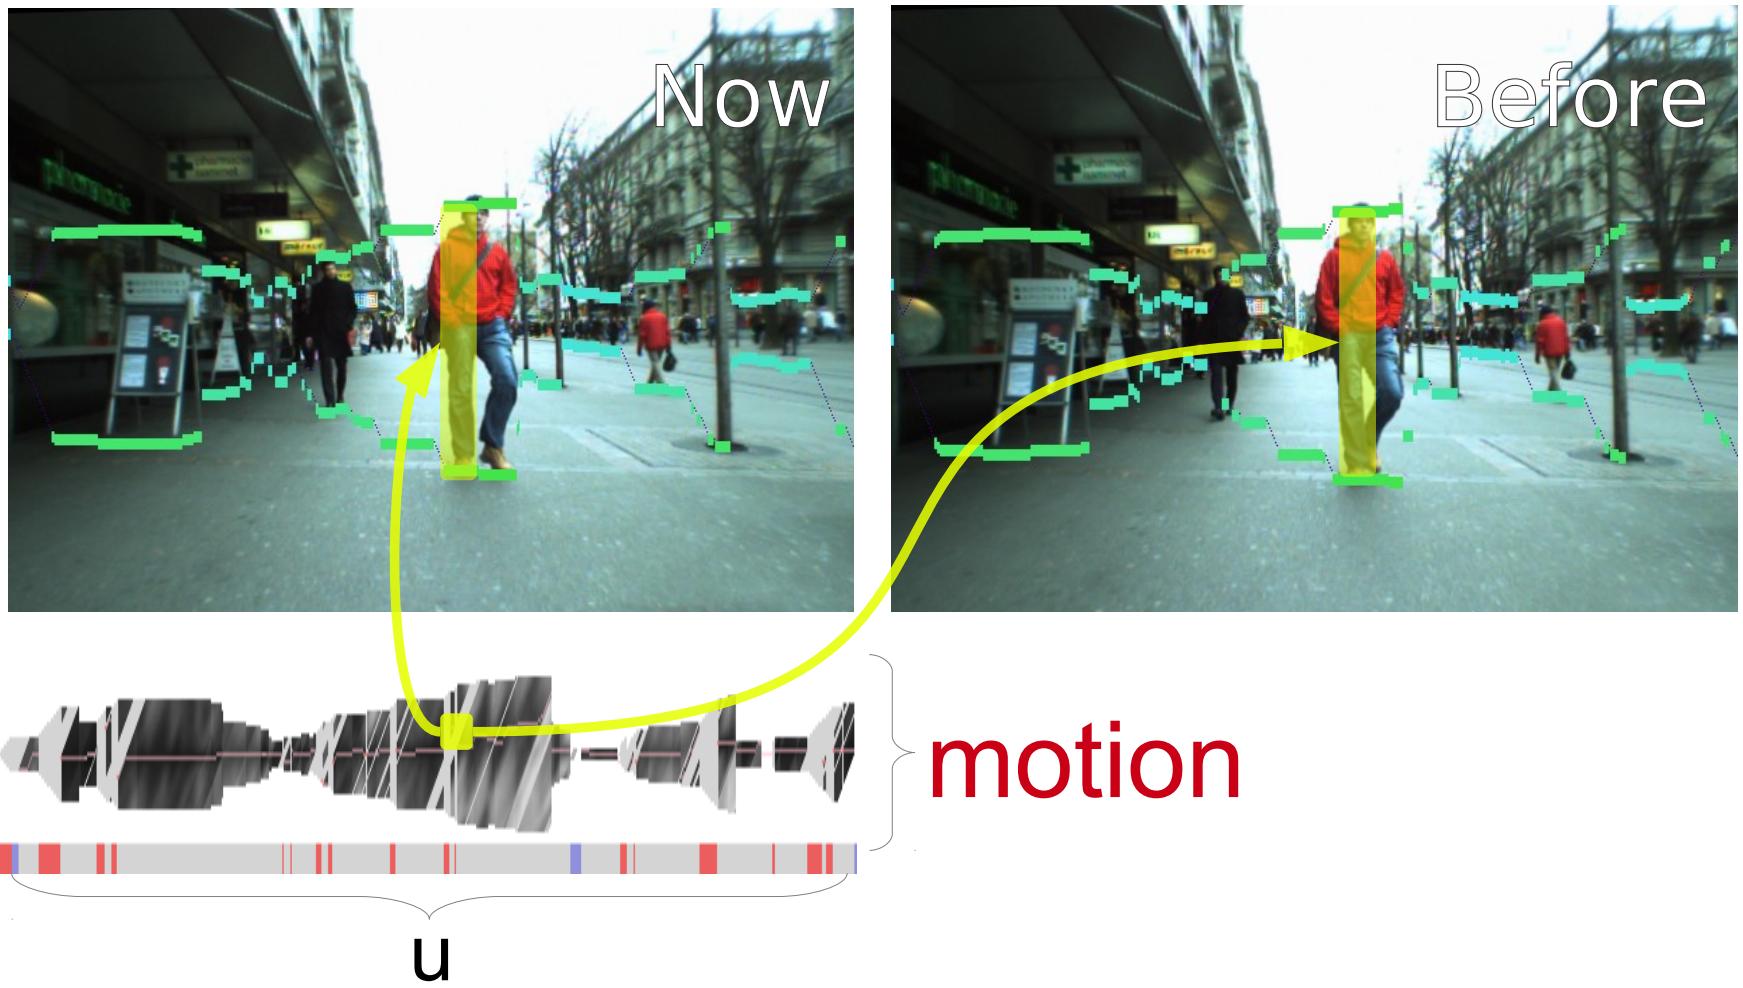
they can be computed very fast  
(300 Hz on CPU, less than 4 ms)



Source code available at  
<http://rodrigob.github.com>

Rodrigo Benenson

# Direct stixels motion estimation



# Direct stixels motion estimation

